

Lab Assignment 2

1. Influence of Regularization

Methods

The dataset I used was called `fashion_mnist`, which can be loaded from `keras.datasets`. It includes 70,000 images in between the training and test sets. I only used the training set to perform my own splits, which includes 60,000 images. Each image is a 28 by 28 grayscale image. The associated labels are made up of 10 classes, each representing a clothing item shown in figure 1.. After using `test_train_split`, I got a training dataset with 42,000 samples and validation and test datasets with 9,000 samples each.

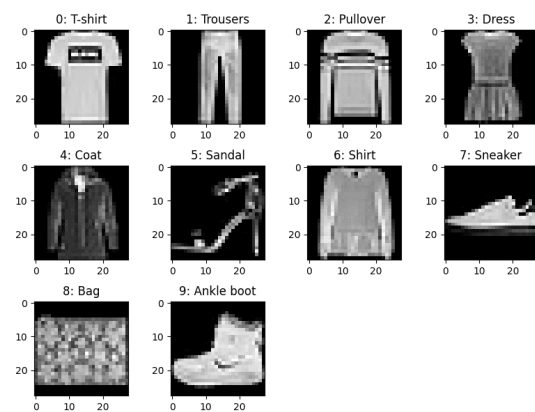


Figure 1 - Dataset Classes

I used `keras` in `tensorflow` to build 8 neural networks. 4 of these models included one convolutional layer and the other 4 included two convolutional layers. For each set of 4, I created one network without using regularization, one network using L2 kernel regularization, one network using a dropout function, and one network using batch normalization.

For these models, there were several hyperparameters that I kept constant. I fit each model using the same training data. In each model, I used mini batch gradient descent with 5 epochs and a batch size of 100. For each convolutional layer, I used a filter size of 16, the `relu` activation function, and a kernel size of 3. After each convolutional layer, I used a pooling layer with a pool size of 2,2. At the end of each model, I used two dense layers. The first dense layer

uses the relu activation function and has a size of 80. The second dense layer uses the softmax activation function and has a size of 10. Each time I used L2 regularization, I used a value of 0.01. This was used on convolutional layers but not on the dense layers. Each time I used dropout, I chose a value of 0.1. Dropout was used after each convolutional layer and after the first dense layer. When Batch Normalization was used, it was done after each convolutional layer and after the first dense layer. Analysis was performed using the validation data.

Results

Of the 8 models I trained. The model that had one convolutional layer and used batch normalization returned the highest validation accuracy. This value was 86.69, shown in the

```
-----
Total training time: 1.346398401260376 minutes.
Training loss: 0.2544
Training accuracy: 91.00
-----
Validation loss: 0.3019
Validation accuracy: 89.69
-----
```




Figure 2 - Optimal Model

Figure 2. The training and validation accuracies are also shown in the plot over each of the 5 epochs. The training time was 1.35 minutes.

After creating each of the 8 neural networks and analyzing the optimal approach, I applied the model on the test data. Instead of using the train data to train the model, I used the training and validation data together.

This resulted in the train / test data split being 85 / 15. I got a final accuracy score of 89.23, shown in

figure 3. The training time was 1.23 minutes.

```
510/510 [=====] - 15s 29ms
-----
Total training time: 1.2284663478533426 minutes.
Test accuracy: 89.23
-----
```

Figure 3 - Final Accuracy

Analysis

One trend observed is that one convolutional layer generally resulted in a higher validation accuracy than two convolutional layers did despite having shorter training times. When no regularization was used, the respective accuracies were 85.41 and 81.24. With L2 regularization, 80.11 and 82.82. With dropout, they were 87.72 and 71.12. Finally, with batch normalization, they were 89.69 and 85.98. My friend found the opposite to be true. Additionally, online sources say that a higher number of convolutional layers usually results in a higher accuracy. I am assuming that I am getting these results because I am using a smaller filter size than my friend. I am using a filter size of 16, while he is using a filter size of 64. With a small filter size, I may be limited in the number of abstractions I can make, and the accuracy may be more likely to decrease for this reason.

Another trend observed is that batch normalization consistently resulted in a higher validation accuracy than the other regularization methods. L2 regularization did not seem to work very well. This method forces the weights towards 0 but does not make them 0. I used a value of 0.01 and it is possible that this value penalized the weights too much. When comparing results from the different number of convolutional layers, both had similar validation accuracies. Over-regularization could have resulted in underfitting in each of these models. When using dropout, I had a good accuracy with one layer but poor accuracy with two layers. This could be explained by dropout having a larger effect a network with a smaller number of neurons. Batch normalization, which reduces the distribution of input variables, was shown to be most effective. It was especially effective with one layer. It is likely that this regularization technique worked the best because it did not have the problems that the other two regularization techniques had. It

is also likely that it was more accurate than no regularization at all because the model containing no regularization had overfitting.

2. Interpreting CNN Representation

Methods

For part b, I visualized the kernel for two different convolutional layers. The assignment asks us to use our model that results in the highest validation accuracy. However, my optimal model only has one convolutional layer. For this reason, I have used a different model to visualize and understand multiple convolutional layers. I will choose the model where I have two convolutional layers and use batch normalization for regularization.

To visualize the convolutional layers, I first selected the convolutional layer and used the `get_weights()` function. From `get_weights`, I extracted the filters and the bias. Next, I printed the shape of each of the filters. The shapes can be shown in each of the images. I created a subplot with the number of columns as the number of inputs and the number of rows as the number of filters in that convolutional layer. I plotted each of the 3 by 3 kernels using a grayscale colormap.

Results


`conv2d_3 (3, 3, 1, 16)`



Figure 4 -1st convolutional layer

For the first convolutional layer, I got 16 different kernels. For the second convolutional layer, I got 256 or $16 * 16$ of these kernels. The kernels for the first layer are shown in figure 4. The kernels for the second layer are shown in figure 5.

conv2d_4 (3, 3, 16, 16)



Analysis

The main difference that can be noted when comparing the first convolutional layer and the second one is that the second layer has way more kernels. The first convolutional layer has a single column with 16 rows. This is because there is only one input, and the filter size is 16. The second convolutional layer has an input size of 16 neurons. When a convolutional layer with a filter size of 16 is applied to this input, we get a 16 by 16 array of 3 by 3 kernels.

Figure 5 - 2nd convolutional layer

One thing that can be inferred from these visualizations is that the abstractions made in the first convolutional layer are important and that it makes sense to have a higher number of filters in the first layer. In the first layer, 16 abstractions are made. Each of these outputs are assigned a weight. In the second layer, there are 256 abstractions and each of these are assigned a weight. The model is simplified if only 16 abstractions are made in the first layer. Because of this, the abstractions are going to be more generalized. As we go further into the network, these abstractions become more specialized. The accuracy does not benefit from this specialization

when the initial layer is too generalized. This would explain why my friend, who used a filter size of 64 was able to benefit more from this specialization in the 2nd layer.

Lab Assignment 2

September 27, 2022

```
[2]: #Reid Glaze  
#Dr Gurari  
#27 Sept 2022  
#Lab Assignment 2  
  
from tensorflow.keras.utils import to_categorical  
from tensorflow.python.keras import Sequential  
from tensorflow.python.keras.layers import Dense, Conv2D, Flatten, MaxPooling2D  
from tensorflow.keras.layers import Dropout  
from tensorflow.keras.layers import BatchNormalization  
  
#import dataset - will only use the train portion to split into train, test, and validation  
from keras.datasets import fashion_mnist  
from keras import regularizers  
import time  
import numpy as np  
from matplotlib import pyplot as plt  
(X, y), (X_test, y_test) = fashion_mnist.load_data()  
  
#train, test, validation split 70-15-15  
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
    random_state=123)  
X_test, X_val, y_test, y_val = train_test_split(X_test, y_test, test_size=0.5,  
    random_state=123)  
  
#Data pre-processing  
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)  
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)  
X_val = X_val.reshape(X_val.shape[0], 28, 28, 1)  
  
y_train = to_categorical(y_train)  
y_test = to_categorical(y_test)  
y_val = to_categorical(y_val)
```

[110]: #1) one convolutional layer, no regularization

```
model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = ["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100, validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = "validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()
```


Model: "sequential_99"

Layer (type)	Output Shape	Param #
conv2d_151 (Conv2D)	(None, 26, 26, 16)	160


```

-----
max_pooling2d_151 (MaxPoolin (None, 13, 13, 16)          0
-----
flatten_99 (Flatten)          (None, 2704)          0
-----
dense_198 (Dense)              (None, 80)          216400
-----
dense_199 (Dense)              (None, 10)          810
=====
Total params: 217,370
Trainable params: 217,370
Non-trainable params: 0
-----
Epoch 1/5
420/420 [=====] - 13s 30ms/step - loss: 1.2219 -
accuracy: 0.7522 - val_loss: 0.4927 - val_accuracy: 0.8254
Epoch 2/5
420/420 [=====] - 12s 29ms/step - loss: 0.4497 -
accuracy: 0.8383 - val_loss: 0.4527 - val_accuracy: 0.8399
Epoch 3/5
420/420 [=====] - 11s 26ms/step - loss: 0.3923 -
accuracy: 0.8575 - val_loss: 0.4150 - val_accuracy: 0.8556
Epoch 4/5
420/420 [=====] - 10s 23ms/step - loss: 0.3530 -
accuracy: 0.8696 - val_loss: 0.4037 - val_accuracy: 0.8587
Epoch 5/5
420/420 [=====] - 9s 23ms/step - loss: 0.3245 -
accuracy: 0.8804 - val_loss: 0.4190 - val_accuracy: 0.8541
-----
Total training time: 0.9325737158457438 minutes.
Training loss: 0.3123
Training accuracy: 88.37
-----
Validation loss: 0.4190
Validation accuracy: 85.41
-----

```



[87]: #2) one convolutional layer, use L2 kernel regularization

```
model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (
    ↪(28,28,1), kernel_regularizer = regularizers.l2(0.01)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = (
    ↪["accuracy"]))

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100, (
    ↪validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
```

```

print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = "validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```

Model: "sequential_84"

Layer (type)	Output Shape	Param #
conv2d_132 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_132 (MaxPoolin	(None, 13, 13, 16)	0
flatten_84 (Flatten)	(None, 2704)	0
dense_168 (Dense)	(None, 80)	216400
dense_169 (Dense)	(None, 10)	810

Total params: 217,370

Trainable params: 217,370

Non-trainable params: 0

Epoch 1/5

420/420 [=====] - 11s 26ms/step - loss: 7.7318 - accuracy: 0.3049 - val_loss: 4.8455 - val_accuracy: 0.6350

Epoch 2/5

420/420 [=====] - 9s 20ms/step - loss: 4.2121 - accuracy: 0.7100 - val_loss: 3.7600 - val_accuracy: 0.7559

Epoch 3/5

420/420 [=====] - 9s 21ms/step - loss: 3.4967 - accuracy: 0.7590 - val_loss: 3.2318 - val_accuracy: 0.7782

Epoch 4/5

420/420 [=====] - 10s 23ms/step - loss: 3.0022 - accuracy: 0.7789 - val_loss: 2.9048 - val_accuracy: 0.7290

Epoch 5/5

420/420 [=====] - 9s 21ms/step - loss: 2.6023 -
accuracy: 0.7902 - val_loss: 2.4195 - val_accuracy: 0.8011


Total training time: 0.7899572650591532 minutes.

Training loss: 2.4131

Training accuracy: 80.23

Validation loss: 2.4195

Validation accuracy: 80.11



[89]: #3) one convolutional layer, use Dropout function

```
model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.1))
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(Dropout(0.1))
model.add(Dense(10, activation = "softmax"))
model.summary()
```

```

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics =
    ↪["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100,
    ↪validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label =
    ↪"validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```


Model: "sequential_86"

Layer (type)	Output Shape	Param #
conv2d_134 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_134 (MaxPoolin	(None, 13, 13, 16)	0
module_wrapper_39 (ModuleWra	(None, 13, 13, 16)	0
flatten_86 (Flatten)	(None, 2704)	0
dense_172 (Dense)	(None, 80)	216400
module_wrapper_40 (ModuleWra	(None, 80)	0
dense_173 (Dense)	(None, 10)	810

```

=====
Total params: 217,370
Trainable params: 217,370
Non-trainable params: 0
-----
Epoch 1/5
420/420 [=====] - 12s 27ms/step - loss: 0.9949 -
accuracy: 0.7462 - val_loss: 0.4719 - val_accuracy: 0.8371
Epoch 2/5
420/420 [=====] - 10s 24ms/step - loss: 0.4566 -
accuracy: 0.8360 - val_loss: 0.3961 - val_accuracy: 0.8632
Epoch 3/5
420/420 [=====] - 10s 25ms/step - loss: 0.4088 -
accuracy: 0.8530 - val_loss: 0.3861 - val_accuracy: 0.8678
Epoch 4/5
420/420 [=====] - 10s 24ms/step - loss: 0.3749 -
accuracy: 0.8663 - val_loss: 0.3584 - val_accuracy: 0.8740
Epoch 5/5
420/420 [=====] - 11s 25ms/step - loss: 0.3537 -
accuracy: 0.8708 - val_loss: 0.3507 - val_accuracy: 0.8772
-----
Total training time: 0.8832791845003763 minutes.
Training loss: 0.2892
Training accuracy: 89.72
-----
Validation loss: 0.3507
Validation accuracy: 87.72
-----

```



[95]: #4) one convolutional layer, batch normalization

```
model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = ["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100, validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
print("-----")
```

```

print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = "validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```

Model: "sequential_92"

Layer (type)	Output Shape	Param #
conv2d_144 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_144 (MaxPoolin	(None, 13, 13, 16)	0
module_wrapper_49 (ModuleWra	(None, 13, 13, 16)	64
flatten_92 (Flatten)	(None, 2704)	0
dense_184 (Dense)	(None, 80)	216400
module_wrapper_50 (ModuleWra	(None, 80)	320
dense_185 (Dense)	(None, 10)	810

Total params: 217,754

Trainable params: 217,562

Non-trainable params: 192

Epoch 1/5

420/420 [=====] - 19s 43ms/step - loss: 0.5430 - accuracy: 0.8140 - val_loss: 0.4056 - val_accuracy: 0.8624


Epoch 2/5

420/420 [=====] - 17s 40ms/step - loss: 0.3714 - accuracy: 0.8726 - val_loss: 0.3607 - val_accuracy: 0.8789


```

Epoch 3/5
420/420 [=====] - 15s 37ms/step - loss: 0.3251 -
accuracy: 0.8876 - val_loss: 0.3247 - val_accuracy: 0.8892
Epoch 4/5
420/420 [=====] - 15s 35ms/step - loss: 0.2949 -
accuracy: 0.8976 - val_loss: 0.3173 - val_accuracy: 0.8953
Epoch 5/5
420/420 [=====] - 14s 34ms/step - loss: 0.2737 -
accuracy: 0.9037 - val_loss: 0.3019 - val_accuracy: 0.8969
-----
Total training time: 1.346398401260376 minutes.
Training loss: 0.2544
Training accuracy: 91.00
-----
Validation loss: 0.3019
Validation accuracy: 89.69
-----

```



[91]: #5) two convolutional layers, no regularization

```

model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPooling2D(pool_size=(2,2)))

```

```

model.add(Conv2D(16, kernel_size = 3, activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = _
↳["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100, _
↳validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = _
↳"validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```

Model: "sequential_88"


Layer (type)	Output Shape	Param #
conv2d_136 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_136 (MaxPoolin	(None, 13, 13, 16)	0
conv2d_137 (Conv2D)	(None, 11, 11, 16)	2320

```

-----
max_pooling2d_137 (MaxPoolin (None, 5, 5, 16)          0
-----
flatten_88 (Flatten)          (None, 400)          0
-----
dense_176 (Dense)             (None, 80)          32080
-----
dense_177 (Dense)             (None, 10)          810
=====
Total params: 35,370
Trainable params: 35,370
Non-trainable params: 0

-----
Epoch 1/5
420/420 [=====] - 13s 31ms/step - loss: 3.6413 -
accuracy: 0.6495 - val_loss: 0.7028 - val_accuracy: 0.7501
Epoch 2/5
420/420 [=====] - 12s 29ms/step - loss: 0.6501 -
accuracy: 0.7604 - val_loss: 0.6127 - val_accuracy: 0.7738
Epoch 3/5
420/420 [=====] - 13s 31ms/step - loss: 0.5825 -
accuracy: 0.7854 - val_loss: 0.5633 - val_accuracy: 0.7959
Epoch 4/5
420/420 [=====] - 15s 36ms/step - loss: 0.5478 -
accuracy: 0.7961 - val_loss: 0.5239 - val_accuracy: 0.8083
Epoch 5/5
420/420 [=====] - 13s 30ms/step - loss: 0.5261 -
accuracy: 0.8030 - val_loss: 0.5141 - val_accuracy: 0.8124
-----
Total training time: 1.1152189135551454 minutes.
Training loss: 0.5029
Training accuracy: 81.42
-----
Validation loss: 0.5141
Validation accuracy: 81.24
-----

```



[92]: #6) two convolutional layers using kernel regularization

```
model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28,28,1),
    kernel_regularizer = regularizers.l2(0.01)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', kernel_regularizer =
    regularizers.l2(0.01)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics =
    ["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100,
    validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
```

```

print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = "validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```

Model: "sequential_89"

Layer (type)	Output Shape	Param #
conv2d_138 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_138 (MaxPoolin	(None, 13, 13, 16)	0
conv2d_139 (Conv2D)	(None, 11, 11, 16)	2320
max_pooling2d_139 (MaxPoolin	(None, 5, 5, 16)	0
flatten_89 (Flatten)	(None, 400)	0
dense_178 (Dense)	(None, 80)	32080
dense_179 (Dense)	(None, 10)	810

Total params: 35,370

Trainable params: 35,370

Non-trainable params: 0

Epoch 1/5

420/420 [=====] - 16s 35ms/step - loss: 1.6282 - accuracy: 0.6485 - val_loss: 0.9356 - val_accuracy: 0.7836


Epoch 2/5

420/420 [=====] - 14s 34ms/step - loss: 0.8677 - accuracy: 0.7877 - val_loss: 0.8074 - val_accuracy: 0.8026

```

Epoch 3/5
420/420 [=====] - 15s 35ms/step - loss: 0.7567 -
accuracy: 0.8105 - val_loss: 0.7263 - val_accuracy: 0.8202
Epoch 4/5
420/420 [=====] - 16s 39ms/step - loss: 0.6794 -
accuracy: 0.8232 - val_loss: 0.6668 - val_accuracy: 0.8278
Epoch 5/5
420/420 [=====] - 16s 37ms/step - loss: 0.6236 -
accuracy: 0.8326 - val_loss: 0.6397 - val_accuracy: 0.8282
-----
Total training time: 1.2868525544802347 minutes.
Training loss: 0.6040
Training accuracy: 83.72
-----
Validation loss: 0.6397
Validation accuracy: 82.82
-----

```



[93]: *#7) two convolutional layers using dropout*

```

model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (
    28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.1))

```

```

model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(Dropout(0.1))
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(Dropout(0.1))
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = ["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100, validation_data= (X_val, y_val))
end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = "validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```


Model: "sequential_90"

Layer (type)	Output Shape	Param #
conv2d_140 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_140 (MaxPoolin	(None, 13, 13, 16)	0

```

-----
module_wrapper_43 (ModuleWra (None, 13, 13, 16)          0
-----
conv2d_141 (Conv2D)          (None, 11, 11, 16)          2320
-----
max_pooling2d_141 (MaxPoolin (None, 5, 5, 16)          0
-----
module_wrapper_44 (ModuleWra (None, 5, 5, 16)          0
-----
flatten_90 (Flatten)          (None, 400)          0
-----
dense_180 (Dense)          (None, 80)          32080
-----
module_wrapper_45 (ModuleWra (None, 80)          0
-----
dense_181 (Dense)          (None, 10)          810
=====
Total params: 35,370
Trainable params: 35,370
Non-trainable params: 0
-----
Epoch 1/5
420/420 [=====] - 18s 41ms/step - loss: 3.7297 -
accuracy: 0.1301 - val_loss: 2.2019 - val_accuracy: 0.2059
Epoch 2/5
420/420 [=====] - 16s 38ms/step - loss: 2.1061 -
accuracy: 0.2911 - val_loss: 1.8523 - val_accuracy: 0.3944
Epoch 3/5
420/420 [=====] - 16s 38ms/step - loss: 1.4171 -
accuracy: 0.4820 - val_loss: 0.9859 - val_accuracy: 0.6600
Epoch 4/5
420/420 [=====] - 16s 38ms/step - loss: 0.9912 -
accuracy: 0.6143 - val_loss: 0.8114 - val_accuracy: 0.7006
Epoch 5/5
420/420 [=====] - 17s 40ms/step - loss: 0.8815 -
accuracy: 0.6559 - val_loss: 0.7465 - val_accuracy: 0.7112
-----
Total training time: 1.3820756196975708 minutes.
Training loss: 0.7578
Training accuracy: 70.34
-----
Validation loss: 0.7465
Validation accuracy: 71.12
-----

```

[94]: #8) two convolutional layer, batch normalization

```
model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28,28,1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())
model.add(Conv2D(16, kernel_size = 3, activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = ["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(X_train, y_train, epochs = 5, batch_size = 100, validation_data= (X_val, y_val))
```

```

end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_train, y_train, verbose = 0)
print("Training loss: %.4f" % score[0])
print("Training accuracy: %.2f" % (score[1] * 100.0))
print("-----")
score = model.evaluate(X_val, y_val, verbose = 0)
print("Validation loss: %.4f" % score[0])
print("Validation accuracy: %.2f" % (score[1] * 100.0))
print("-----")
scores.append(score[1])
histories.append(history)
plt.plot(histories[0].history["accuracy"], color = "blue", label = "train")
plt.plot(histories[0].history["val_accuracy"], color = "red", label = "validation")
plt.ylabel("Accuracy")
plt.xlabel("Num of epochs")
leg = plt.legend(loc='lower center')
plt.show()

```

Model: "sequential_91"

Layer (type)	Output Shape	Param #
conv2d_142 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_142 (MaxPoolin	(None, 13, 13, 16)	0
module_wrapper_46 (ModuleWra	(None, 13, 13, 16)	64
conv2d_143 (Conv2D)	(None, 11, 11, 16)	2320
max_pooling2d_143 (MaxPoolin	(None, 5, 5, 16)	0
module_wrapper_47 (ModuleWra	(None, 5, 5, 16)	64
flatten_91 (Flatten)	(None, 400)	0
dense_182 (Dense)	(None, 80)	32080
module_wrapper_48 (ModuleWra	(None, 80)	320
dense_183 (Dense)	(None, 10)	810

Total params: 35,818

Trainable params: 35,594

Non-trainable params: 224

```
-----  
Epoch 1/5  
420/420 [=====] - 29s 66ms/step - loss: 0.7374 -  
accuracy: 0.7545 - val_loss: 0.5230 - val_accuracy: 0.8198  
Epoch 2/5  
420/420 [=====] - 25s 60ms/step - loss: 0.4755 -  
accuracy: 0.8340 - val_loss: 0.4464 - val_accuracy: 0.8452  
Epoch 3/5  
420/420 [=====] - 24s 57ms/step - loss: 0.4170 -  
accuracy: 0.8523 - val_loss: 0.4090 - val_accuracy: 0.8600  
Epoch 4/5  
420/420 [=====] - 25s 59ms/step - loss: 0.3846 -  
accuracy: 0.8639 - val_loss: 0.3899 - val_accuracy: 0.8677  
Epoch 5/5  
420/420 [=====] - 24s 57ms/step - loss: 0.3630 -  
accuracy: 0.8714 - val_loss: 0.3979 - val_accuracy: 0.8598  
-----
```

Total training time: 2.116146699587504 minutes.

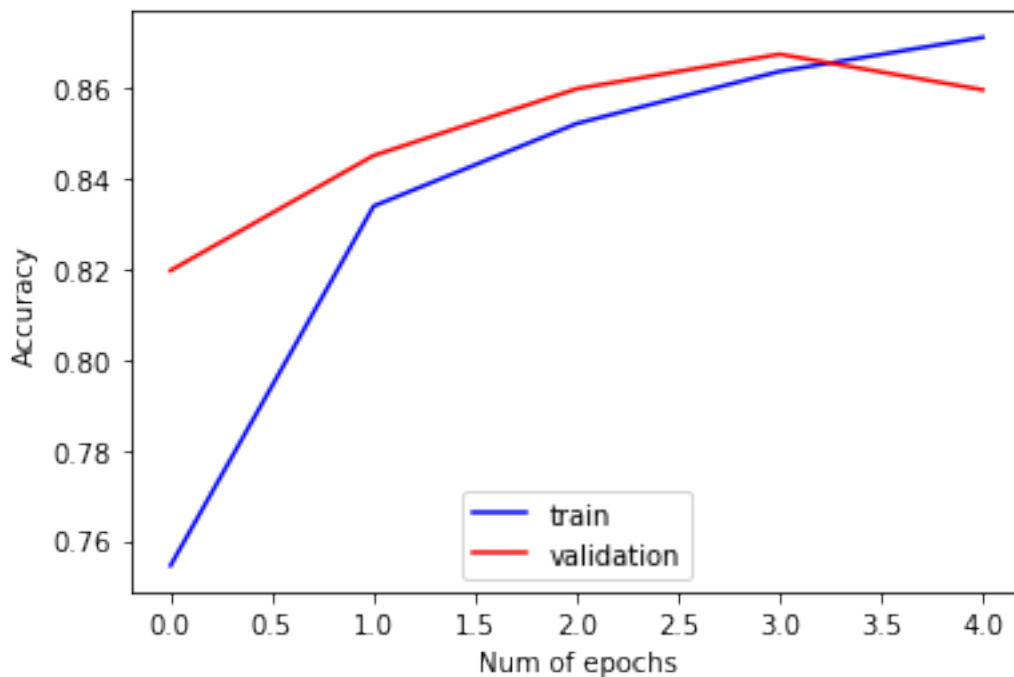
Training loss: 0.3704

Training accuracy: 86.47

```
-----  
Validation loss: 0.3979
```

Validation accuracy: 85.98

```
-----
```



```
[10]: #Train best model with train & validation data and measure accuracy using test
      ↪ data.
      #combine train and validation data
      train_data = list(X_train)
      train_labels = list(y_train)
      val_data = list(X_val)
      val_labels = list(y_val)
      all_data = train_data+val_data
      all_labels = train_labels+val_labels
      new_X_train = np.array(all_data)
      new_y_train = np.array(all_labels)

      model = Sequential()
      model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape =
      ↪ (28,28,1)))
      model.add(MaxPooling2D(pool_size=(2,2)))
      model.add(BatchNormalization())
      model.add(Flatten())
      model.add(Dense(80, activation = 'relu'))
      model.add(BatchNormalization())
      model.add(Dense(10, activation = "softmax"))
      model.summary()

      model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics =
      ↪ ["accuracy"])

      scores = []
      histories = []
      start = time.time()
      history = model.fit(new_X_train, new_y_train, epochs = 5, batch_size = 100)
      end = time.time()
      num_mins = (end-start)/60
      print("-----")
      print("Total training time: " + str(num_mins) + " minutes.")
      score = model.evaluate(X_test, y_test, verbose = 0)
      print("Test accuracy: %.2f" % (score[1] * 100.0))
      print("-----")
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_2 (MaxPooling2D)	(None, 13, 13, 16)	0
module_wrapper_4 (ModuleWrap)	(None, 13, 13, 16)	64

```

-----
flatten_2 (Flatten)                (None, 2704)                0
-----
dense_4 (Dense)                    (None, 80)                   216400
-----
module_wrapper_5 (ModuleWrap (None, 80)                320
-----
dense_5 (Dense)                    (None, 10)                   810
=====
Total params: 217,754
Trainable params: 217,562
Non-trainable params: 192
-----
Epoch 1/5
510/510 [=====] - 16s 28ms/step - loss: 0.5134 -
accuracy: 0.8281
Epoch 2/5
510/510 [=====] - 14s 27ms/step - loss: 0.3517 -
accuracy: 0.8775
Epoch 3/5
510/510 [=====] - 15s 29ms/step - loss: 0.3106 -
accuracy: 0.8928
Epoch 4/5
510/510 [=====] - 14s 28ms/step - loss: 0.2831 -
accuracy: 0.9001
Epoch 5/5
510/510 [=====] - 15s 29ms/step - loss: 0.2652 -
accuracy: 0.9071
-----
Total training time: 1.2284663478533426 minutes.
Test accuracy: 89.23
-----

```

```

[14]: import tensorflow as tf
import glob
import numpy as np
import pickle
import os

from matplotlib import pyplot as plt
from PIL import Image
from urllib import request
from io import BytesIO

from keras.applications.vgg16 import preprocess_input
from keras.preprocessing import image
from keras.applications.vgg16 import VGG16

```

```
[61]: layer = model.layers[0]
filters, bias = layer.get_weights()
print(layer.name, filters.shape)

f_min, f_max = filters.min(), filters.max()
filters = (filters-f_min) / (f_max - f_min)

n_filters = 16; ix = 1
for i in range(n_filters):
    f = filters[:, :, :, i]

    ax = plt.subplot(n_filters, 1, ix)
    ax.set_xticks([])
    ax.set_yticks([])
    plt.imshow(f[:, :], cmap = "gray")
    ix +=1

plt.show()
```

conv2d_3 (3, 3, 1, 16)



```
[48]: model = Sequential()
model.add(Conv2D(16, kernel_size = 3, activation = 'relu', input_shape = (28, 28, 1)))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())
```

```

model.add(Conv2D(16, kernel_size = 3, activation = 'relu'))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())
model.add(Flatten())
model.add(Dense(80, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(10, activation = "softmax"))
model.summary()

model.compile(loss = "categorical_crossentropy", optimizer = "SGD", metrics = _
↳ ["accuracy"])

scores = []
histories = []
start = time.time()
history = model.fit(new_X_train, new_y_train, epochs = 5, batch_size = 100)
end = time.time()
num_mins = (end-start)/60
print("-----")
print("Total training time: " + str(num_mins) + " minutes.")
score = model.evaluate(X_test, y_test, verbose = 0)
print("Test accuracy: %.2f" % (score[1] * 100.0))
print("-----")

```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_3 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_3 (MaxPooling2)	(None, 13, 13, 16)	0
module_wrapper_6 (ModuleWrap	(None, 13, 13, 16)	64
conv2d_4 (Conv2D)	(None, 11, 11, 16)	2320
max_pooling2d_4 (MaxPooling2)	(None, 5, 5, 16)	0
module_wrapper_7 (ModuleWrap	(None, 5, 5, 16)	64
flatten_3 (Flatten)	(None, 400)	0
dense_6 (Dense)	(None, 80)	32080
module_wrapper_8 (ModuleWrap	(None, 80)	320
dense_7 (Dense)	(None, 10)	810

Total params: 35,818
Trainable params: 35,594
Non-trainable params: 224

```
-----  
Epoch 1/5  
510/510 [=====] - 27s 51ms/step - loss: 0.6815 -  
accuracy: 0.7699  
Epoch 2/5  
510/510 [=====] - 22s 43ms/step - loss: 0.4425 -  
accuracy: 0.8439  
Epoch 3/5  
510/510 [=====] - 20s 40ms/step - loss: 0.3895 -  
accuracy: 0.8616  
Epoch 4/5  
510/510 [=====] - 20s 40ms/step - loss: 0.3602 -  
accuracy: 0.8713  
Epoch 5/5  
510/510 [=====] - 20s 40ms/step - loss: 0.3402 -  
accuracy: 0.8795  
-----
```

Total training time: 1.8317898988723755 minutes.
Test accuracy: 86.64

```
[72]: layer = model.layers[0]  
      filters, bias = layer.get_weights()  
      print(layer.name, filters.shape)  
  
      n_filters = 16; ix = 1  
      for i in range(n_filters):  
          f = filters[:, :, :, i]  
  
          ax = plt.subplot(n_filters, 1, ix)  
          ax.set_xticks([])  
          ax.set_yticks([])  
          plt.imshow(f[:, :, :], cmap = "gray")  
          ix += 1  
  
      plt.show()  
  
      layer = model.layers[0]  
      filters, bias = layer.get_weights()  
  
      layer = model.layers[3]
```



```

filters, bias = layer.get_weights()
print(layer.name, filters.shape)

n_filters = 16; ix = 1
for i in range(n_filters):
    f = filters[:, :, :, i]

    #plot each channel
    for j in range(16):
        ax = plt.subplot(n_filters, 16, ix)
        ax.set_xticks([])
        ax.set_yticks([])
        plt.imshow(f[:, :, j], cmap = "gray")
        ix +=1

plt.show()

```

conv2d_3 (3, 3, 1, 16)



conv2d_4 (3, 3, 16, 16)



[]: